IBM Cloud

Introduction to Docker

Hands-on Workshop

Lab Guide

Notices and Disclaimers

The information contained in these materials is provided for informational purposes only, and is provided AS IS without warranty of any kind, express or implied. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, these materials. Nothing contained in these materials is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software. References in these materials to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. This information is based on current IBM product plans and strategy, which are subject to change by IBM without notice. Product release dates and/or capabilities referenced in these materials may change at any time at IBM's sole discretion based on market opportunities or other factors, and are not intended to be a commitment to future product or feature availability in any way.

This document is current as of the initial date of publication and may be changed by IBM at any time. Not all offerings are available in every country in which IBM operates.

IBM, the IBM logo and ibm.com are trademarks or registered trademarks of International Business Machines Corporation in the United States, other countries, or both. If these and other IBM trademarked terms are marked on their first occurrence in this information with a trademark symbol (® or ™), these symbols indicate U.S. registered or common law trademarks owned by IBM at the time this information was published. Such trademarks may also be registered or common law trademarks in other countries. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at ibm.com/legal/copytrade.shtml

Other company, product and service names may be trademarks or service marks of others

Document Revision History

Rev #	File Name	Date
1.0	Docker and Kubernetes - Lab	1/21/2018

Prepared & Revised by:
Louis V. Frolio - louis.frolio@ibm.com
David Solomon - dlsolomo@us.ibm.com

Table of Contents

Lab Environment Overview

Installed Software and Tools

Software	Link
Docker	https://www.docker.com
VirtualBox	https://www.virtualbox.org/wiki/Downloads
Minikube	https://kubernetes.io/docs/getting-started-guides/minikube/
Docker Hub	https://hub.docker.com/

Section 1: Pre-requisites

Purpose:	This section introduces Docker editions, and installation types along with guidance on how to install Docker on its supported platforms. You will determine the version of Docker running on your laptop and then verify that Docker is running without issues.

Tasks:	Docker software for: • Discuss Docker Editions • Review supported platforms • Gather information about Docker installation • Verify Docker is running without issue

Section 1: Lab Workflow Overview

1. Docker Editions
2. Installation Types
3. Platform Installations
4. Determine Docker Version
5. Verify Docker Installation

Section 1: Lab Instructions

Step	Action
1	**Docker Editions** • Community Edition (CE) Free, quarterly release cadence, no premium support (community). **This lab uses Docker CE** • Enterprise Edition (EE) Not free, quarterly release cadence, premium support available Certified on specific platforms Extra products • General Availability (GA) vs Beta (Edge) **GA** is the stable production ready release of Docker; quarterly cadence for CE and EE **Edge** is the monthly beta release, gets new functionality first. Each edge release is supported for the month, once new edge release is released prior edge cannot get support. Aggregate edge functionally changes roll up into EE GA release quarterly.
2	**Installation Types** • **Direct** Direct installation on a supported operating system. E.g. Linux, RaspPI, Mainframe, Windows Server 2016 • **Mac/Windows 10** Natively does not support “direct” installation of Docker. Small VM (transparent) is spun up to run the containers in. • **Cloud** IBM Cloud, AWS, Azure, Google Cloud Usually have proprietary features specific to cloud
3	**Platform Installations** • **Mac** Install Docker for Mac. For older Mac’s with less than OSX Yosemite 10.10.3 install the Docker Toolbox • **Windows 10 Pro/Enterprise** Install "Docker for Windows" from the Docker Store • **Windows 7, 8, 10 Home** Install Docker Toolbox. Lack of Hyper-V necessitates this type of installation

Step	Action
4	**Determine Docker Version** a. Open terminal (MAC), Shell (Windows), or Quickstart Terminal (Docker Toolbox) then type: ~$ docker version 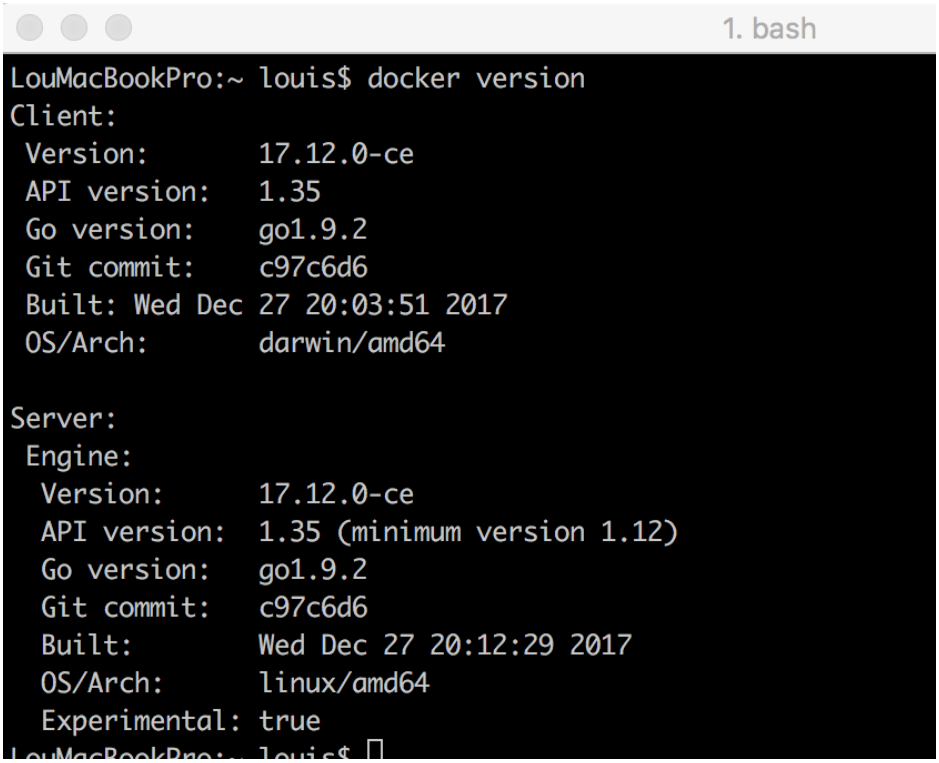 b. You output should be similar.
5	**Verify Docker Installation** a. Run test Docker container ~$ docker container run hello-world

Step	Action
	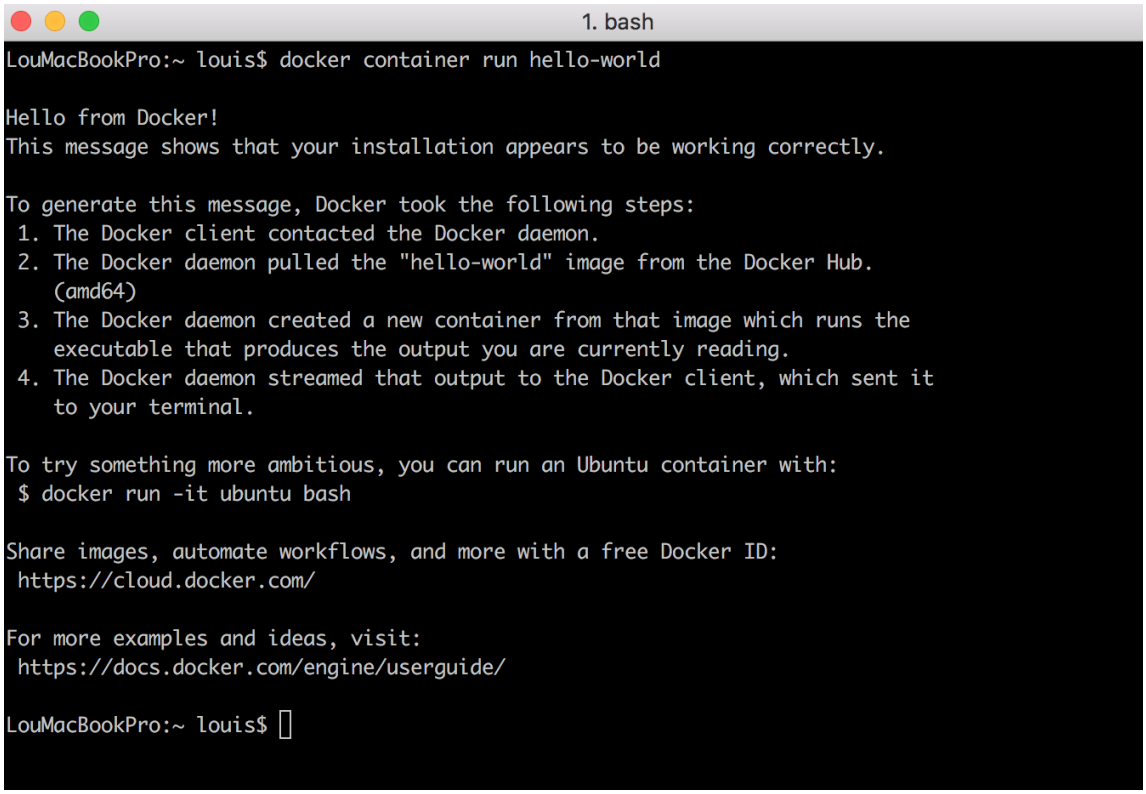 Output verifies that Docker is running and you are able to pull images from Docker Hub, and then start a container from the image.

Section 1: Lab Summary

In this lab, you learned about Docker Editions, installation types, and various platforms supported by Docker. You also ran Docker commands to determine the version of Docker, and to verify that the installation was successful.

Other useful Docker commands:

~$ docker Info - display Docker system-wide information
~$ docker help – display help topics available

Section 2: Container Basics

Purpose:	This lab introduces container basics. You will learn how to run, inspect and manage multiple containers. Also, you will work through establishing console access within the container.

Tasks:	Tasks you will complete in this lab exercise include: • Running containers • Inspecting containers • Container process monitoring • Container shell access

Section 2: Lab Workflow Overview

1. Run a Container
2. Stop/Delete a Container
3. Inspect a Running Container
4. Run Shell Inside a Container

Section 2: Lab Instructions

Step	Action
1	**Run a Container** a. In a terminal window type the following: ~$ docker container run --name webserver -p 8080:80 httpd 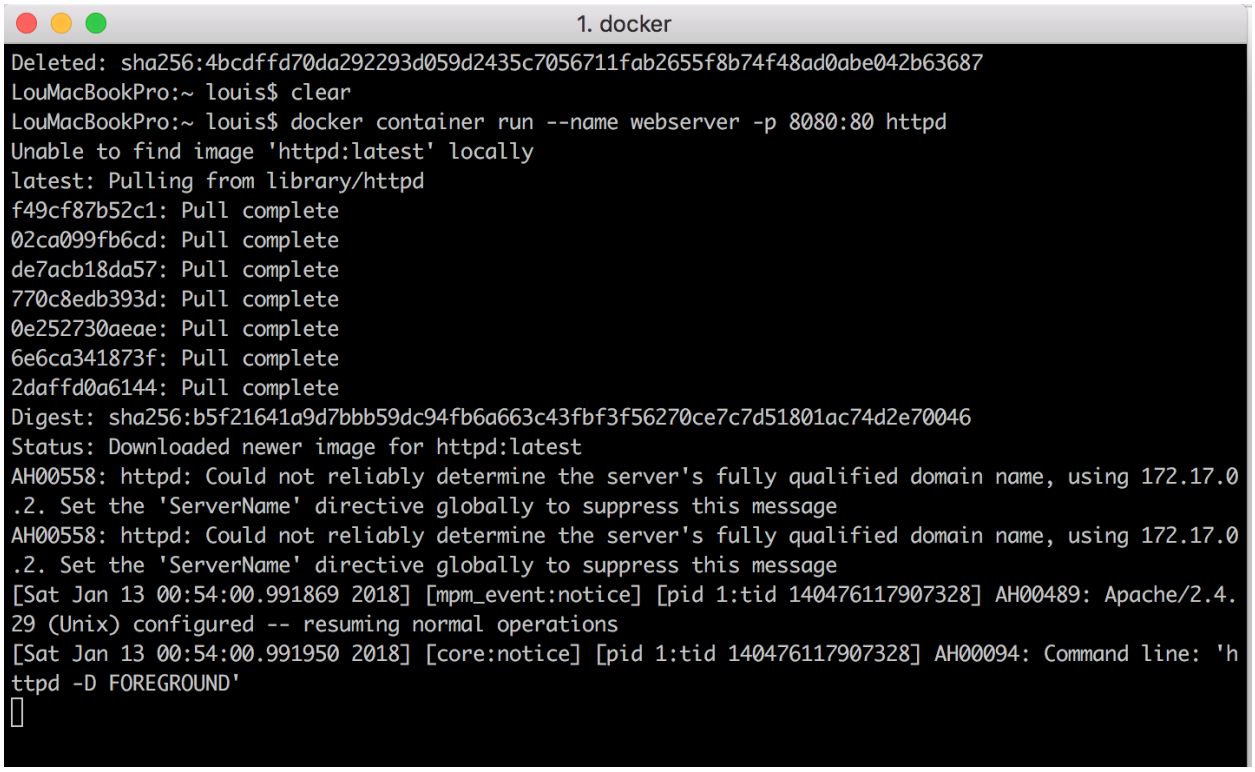 The container “webserver” is an instance of the image “httpd” running as a process. The image is pulled (first time) from the default Docker registry called Docker Hub. There is no limit to the number of containers that can be run from an image. Commands: **--name** – Specify a unique name for the container service. If omitted Docker will create a random, human readable name. **-p** – Specify that the container internal port (80) be exposed to port 8080 on the host. b. Open a browser and navigate to: localhost:8080. You should be presented with a message returned from the web browser “It Works”

Step	Action
	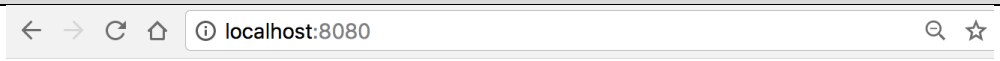 **It works!**
2	**Stop/Delete a Container** a. You stop the container by typing cntrl-c ~$ <Cntrl-c> b. Verify that the container is no longer running: ~$ docker container ps c. Although the container is not running it still exists: ~$ docker container ps -a 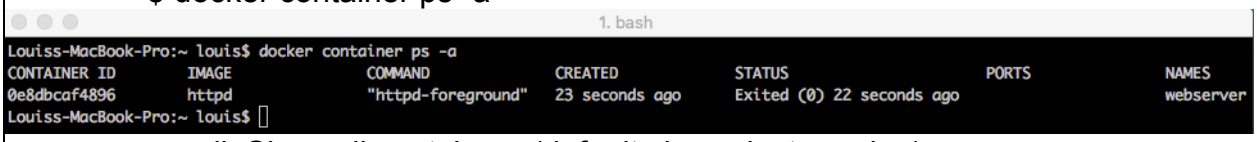 -a, --all: Show all containers (default shows just running) d. Remove the container: ~$ docker container rm webserver Containers can be removed either by their name or container id
3	**Inspect a Running Container** a. Run a new Docker container: ~$ docker run --publish 80:80 --detach --name webserver nginx You should be brought back to the terminal prompt.

Step	Action
	b. Open a browser and navigate to “localhost”. You should be prompted with “Welcome to nginx!” c. Inspect the log file for the running container “webserver”: ~$ docker container logs webserver Click refresh on the browser then re-run the “docker logs …” command. Notice that there is a new log entry for the event. d. Examine the processes running in the container: ~$ docker container top webserver The output shows two nginx processes running in the container. One is the master process, the other is a child process 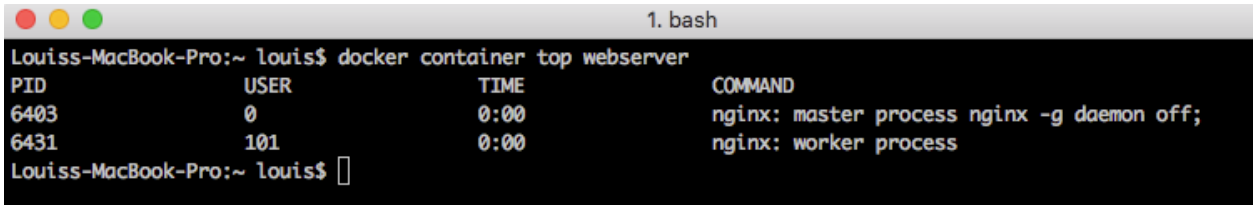 e. Inspect meta-data for running container: ~$ docker container inspect webserver f. Stream live performance container metrics: ~$ docker container stats webserver g. Clean up ~$ docker container rm -f webserver Commands: **-d, --detach** - Run the container in the background.
4	**Run Shell Inside a Container** a. Run container interactively: ~$ docker container run -it --name linuxlight alpine Alpine is a lightweight linux distribution

Step	Action
	b. Run Linux commands in container: # ls -tal // List directories and files # df -h // List file systems and their usage # apk add --no-cache curl // Install curl # curl localhost # exit // Exit shel and stop container c. Start an existing (stopped) container & run Linux commands: ~$ docker container start -ai linuxlight # curl localhost # exit d. Access an already running container: ~$ docker container run -p 80:80 -d --name webserver nginx ~$ docker container ls -a ~$ docker container exec -it webserver bash # ls -tal # exit ~$ docker container ls -a Commands: -i - Run interactively -t - Create pseudo tty -a - Attach to STDIN, STDOUT or STDERR exec - Run a command in a running container run - Run a command in a new container

Section 2: Lab Summary

In this section you learned how to create new containers based on images stored in Docker Hub. You also learned how to interact with containers both from the outside (top, inspect, stats, …), and from the inside (docker exec and run). Access to the Docker service via tty was demonstrated and you learned how to run Linux commands inside the container just as if you were working with a Linux OS.

Section 3: Container Networking Basics

Purpose:	In this lab you will learn basic skills as it relates to networking (security & DNS) with Docker containers:

Tasks:	Tasks you will complete in this lab exercise include: • Determine active ports • Identify container ip address • Create virtual network • Attach container to virtual network

Section 3: Lab Workflow Overview

1. Container Network Attributes
2. Local Network Attributes
3. Virtual Networking
4. Domain Name System (DNS)

Section 3: Lab Instructions

Step	Action
1	**Container Network Attributes** a. On what port is container listening: ~$ docker container port webserver 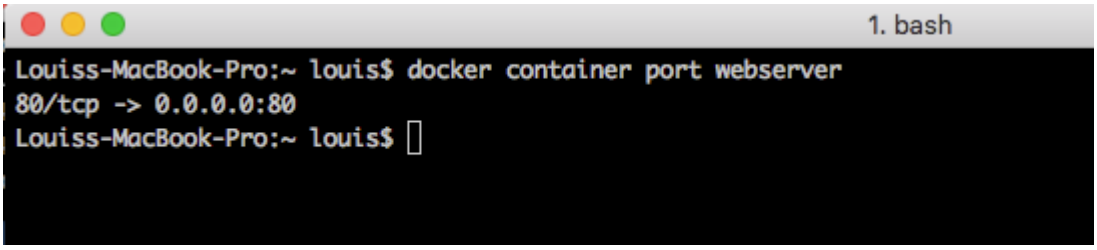 b. Determine ip and mac address of container (Go templating): ~$ docker container inspect --format '{{.NetworkSettings.IPAddress }} | {{.NetworkSettings.MacAddress}}' webserver 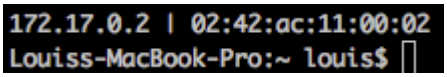
2	**Local Network Attributes** a. List all networks available to Docker: ~$ docker network ls Network named “bridge” is the default network available to Docker b. Get details for local network “bridge”: ~$ docker network inspect bridge Shows containers attached to the “bridge” network. “bridge” is the default network that bridges through NAT firewall to the physical network to which the host is connected.

Step	Action
	 c. Fill in
3	**<u>Virtual Networking</u>** a. Create Virtual Network:

Step	Action
	~$ docker network create mynetwork ~$ docker network ls b. Create new container and run it on “mynetwork”: ~$ docker container run -d --name webserver2 --network mynetwork nginx c. Verify new container is running on “mynetwork”: ~$ docker network inspect mynetwork d. Connect container from “bridge” to “mynetwork”: ~$ docker container inspect webserver | tail -30 ~$ docker network connect mynetwork webserver e. Ajfd;lak
4	**Domain Name System (DNS)** In a Docker environment you cannot have two containers with the same name. Further, because of the dynamic nature of Docker you cannot assign (nor would you want to) a static IP to a container. Instead, you rely on unique container names and reference those in your environment. This is what we have been doing throughout this tutorial. However, you may find that you want to load balance a workload across several containers, but you don’t want the hassle of having to choose. In this case Docker provides the ability to create a network alias that will permit two or more containers to be referenced by a network alias. a. Create a new virtual network called “icecream”: ~$ docker network create icecream ~$ docker network ls b. Create two Elasticsearch containers on the network “icecream”. We don’t specify a port for these, we will work within the virtual network. ~$ docker container run -d --net icecream --net-alias search elasticsearch:2 ~$ docker container run -d --net icecream --net-alias search elasticsearch:2 c. Run a DNS lookup to verify the two containers are on the same network, and share the same DNS name. ~$ docker container run --rm --net icecream alpine nslookup search

Step	Action
	d. Test that DNS round robin is working: ~$ docker container run --rm --net icecream centos curl -s search:9200 Run this command several times and you will see that the Elasticsearch server name will change. Due to DNS caching you may get the same server a few times in a row. e. Clean up: ~$ docker container rm -f <cont1> <cont2>

Section 3: Lab Summary

In this section you were introduce to virtual networking in Docker and how to create and manage them. Also, you learned how to create network aliases and how they can be inspected and used for load balancing

Section 4: Data Persistence in Docker

Purpose:	In this section you will learn the different ways in which Docker can handle and manage data persistence.

Tasks:	Tasks you will complete in this lab exercise include: • Create Docker volumes • Create Bind mounts • Work with volumes and bind mounts in containers

Section 4: Lab Workflow Overview

1. Docker Volumes
2. Docker Bind Mounts

Section 4: Lab Instructions

Step	Action
1	**Docker Volumes** a. Create MySQL container: ~$ docker container run -d --name mysql -e MYSQL_ALLOW_EMPTY_PASSWORD=True mysql b. Open bash shell on container create database objects: ~$ docker container exec -it mysql bash # mysql -u root mysql> show databases; mysql> create database apple; mysql> exit; # exit c. Stop, then restart container to verify that db changes still exist: ~$ docker container stop mysql ~$ docker container start mysql ~$ docker container exec -it mysql bash # mysql -u root mysql> show databases; mysql> exit; # exit ~$ docker container rm -f mysql d. Create a MySQL database and specify volume creation: ~$ docker container run -d --name mysql -e MYSQL_ALLOW_EMPTY_PASSWORD=True -v mysql-db:/var/lib/mysql mysql ~$ docker container ls ~$ docker volume ls ~$ docker volume inspect mysql-db ~$ docker container inspect mysql ** Note that the command above is wrapped. There are no CRLF's. ** The path specified is for a Mac, Windows will be different e. Open bash shell on container create database objects: ~$ docker container exec -it mysql bash # mysql -u root

Step	Action
	mysql> show databases; mysql> create database apple; mysql> exit; # exit f. Delete container “mysql” then recreate and verify that db changes remain: ~$ docker container rm -f mysql ~$ docker volume ls ~$ docker container run -d --name mysql3 -e MYSQL_ALLOW_EMPTY_PASSWORD=True -v mysql-db:/var/lib/mysql mysql ~$ docker container exec -it mysql3 bash # mysql -u root mysql> show databases; mysql> create database apple; mysql> exit; # exit g. Clean up ~$ docker container rm -f mysql3 ~$ docker volume rm mysql-db ~$ docker volume ls ~$ docker container ls -a Volumes are extremely useful for local development projects. You can maintain several volumes to which you can attach a new database that fits a specific purpose.
2	**Docker Bind Mounts** Bind mounts differ from volumes in that they are mappings from the host’s file or directory into a container file or directory. Like volumes, when you delete a container the data is not lost. a. Pull down in index.html from Github ~$ wget https://github.com/team-wolfpack/Docker-and-Kubernetes-Hands-On/blob/master/index.html Inspect the contents of the html file.

Step	Action
	b. Create nginx container and bind mount local directory to specified path within container: ~$ docker container run -d --name nginx -p 80:80 -v $(pwd):/usr/share/nginx/html nginx c. Verify that the correct html is being used: ~$ curl localhost You should see body text of the html file. d. Get to shell prompt in container and verify bind mount mapping ~$ docker container exec -it mysql bash # cd /user/share/nginx/html # cat index.html e. Open another terminal window on the host and change the index.html file. ~$ docker container exec -it mysql bash # cd /user/share/nginx/html # cat index.html

Section 4: Lab Summary

In this lab you were introduced to two approaches to persist data on the host file system. With volumes the container references a volume object on the local file system. Bind mounts create a common link between the local file system and the file system in the container. A bound mount can reference either a file or directory on the host file system.

Section 5: Getting Started with Minikube

Purpose:	In this lab you will learn basic skills as it relates to networking (security & DNS) with Docker containers:

Tasks:	Tasks you will complete in this lab exercise include: • Start Minikube • Configure Kubernetes • Launch & Explore Kubernetes Dashboard

Section 5: Getting Started with Minikube

1. Open Terminal / Launch Minikube
2. Configure Minikube Environment
3. Start Kubernetes Cluster
4. Explore Kubernetes Dashboard

Section 5: Lab Instructions

Step	Action
1	**Start Kubernetes** a. Open a terminal window and type the following command. This will start the Kubernetes cluster. ~$ minikube start 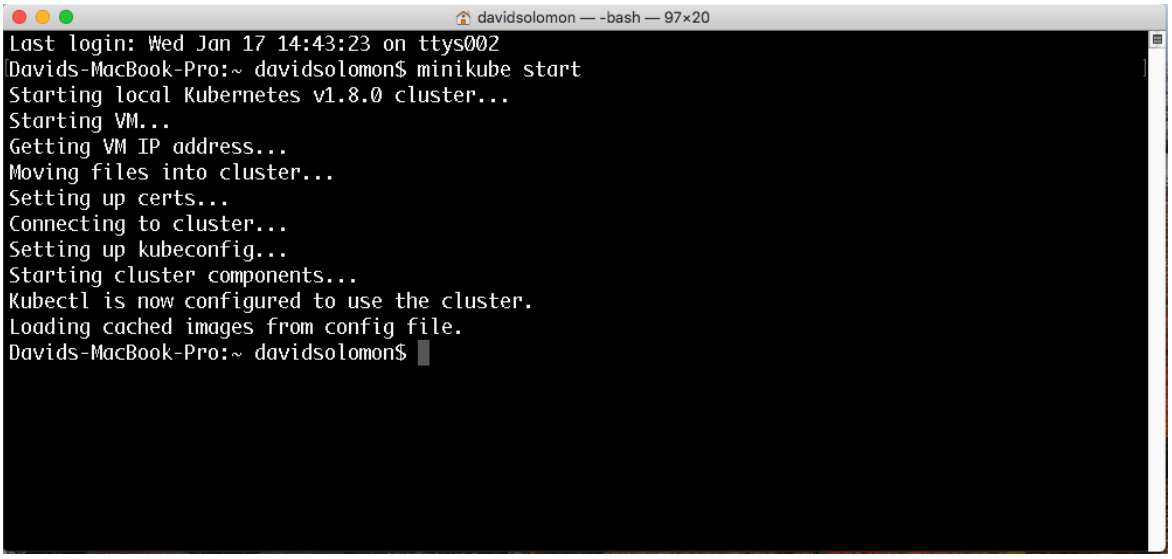 b. Confirm that Minikube has successfully started, as shown above.
2	**Configure the Kubernetes Environment** a. Run the following command to set the current Kubernetes environment to our local cluster (the default cluster name is “minikube”) ~$ kubectl config use-context minikube b. Confirm the output is as shown below:

Step	Action
	
3	**Start the Kubernetes Dashboard** a. Start the Kubernetes web dashboard by entering the following command: ~$ minikube dashboard b. This will open the dashboard in your default browser. You will see in the dashboard the status of any existing deployments, PODS, and services, as shown below: 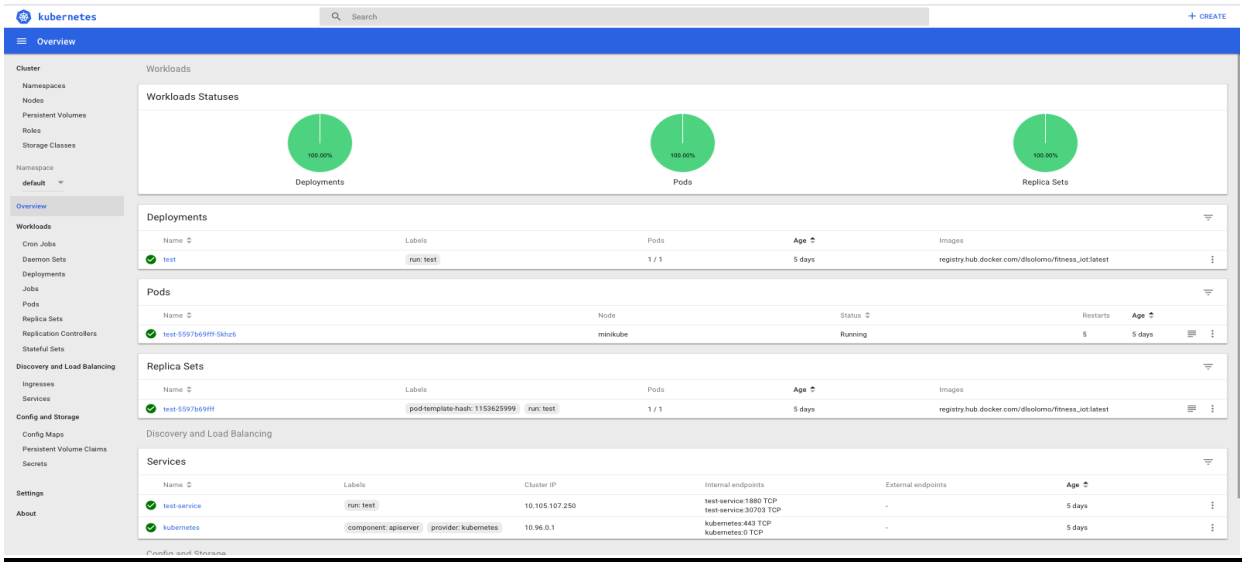

Section 5: Lab Summary

In this section, you learned how to start a Kubernetes cluster on your local machine and viewed the Kubernetes dashboard.

Section 6: Deploy an Application to Kubernetes

Purpose:	In this lab you will learn how to deploy an application to Kubernetes.

Tasks:	Tasks you will complete in this lab exercise include: • Deploy a Docker application to Kubernetes • Expose the application through a service • Access the running application

Section 6: Deploy an Application to Kubernetes

1. Deploy a Docker application to Kubernetes
2. Expose Application through Service
3. Access the Running Application

Section 6: Lab Instructions

Step	Action
1	**Deploy a Docker application to the Kubernetes cluster** a. We will now deploy the same Docker application to your cluster. To do this, enter the following command to create a new deployment called “webserver”: ~$ kubectl run webserver --image=httpd --port=80 b. Confirm the output is as shown below: 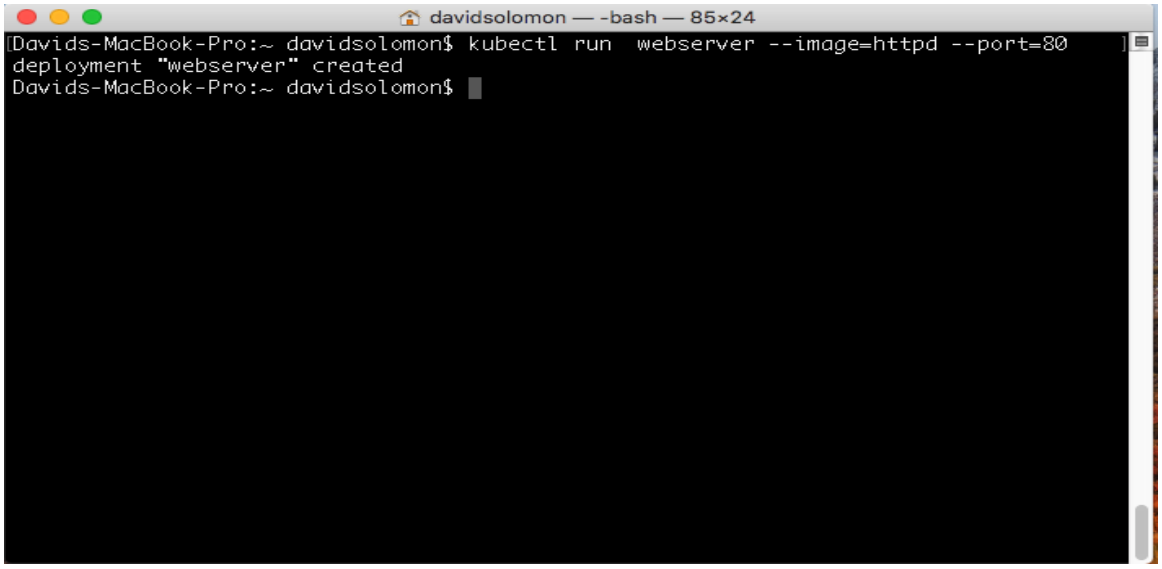 c. Return to the Kubernetes dashboard. You will see that a new deployment and Pod have been created for the application, as shown below, indicating that the application is now running in the cluster. 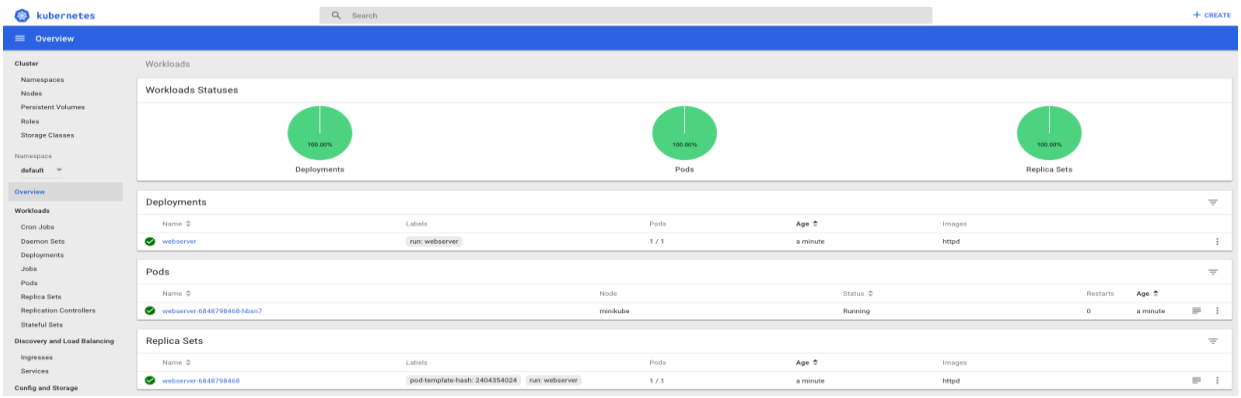

Step	Action
2	**Exposing the application through a service** a. In order to interact with your application from outside the cluster, you will need to create a service which provide an endpoint to expose the application. To do this, enter the following command to create a new service called “webservice”: ~$ kubectl expose deployment webserver --type=NodePort --name *webservice* b. Confirm the output is as shown below: 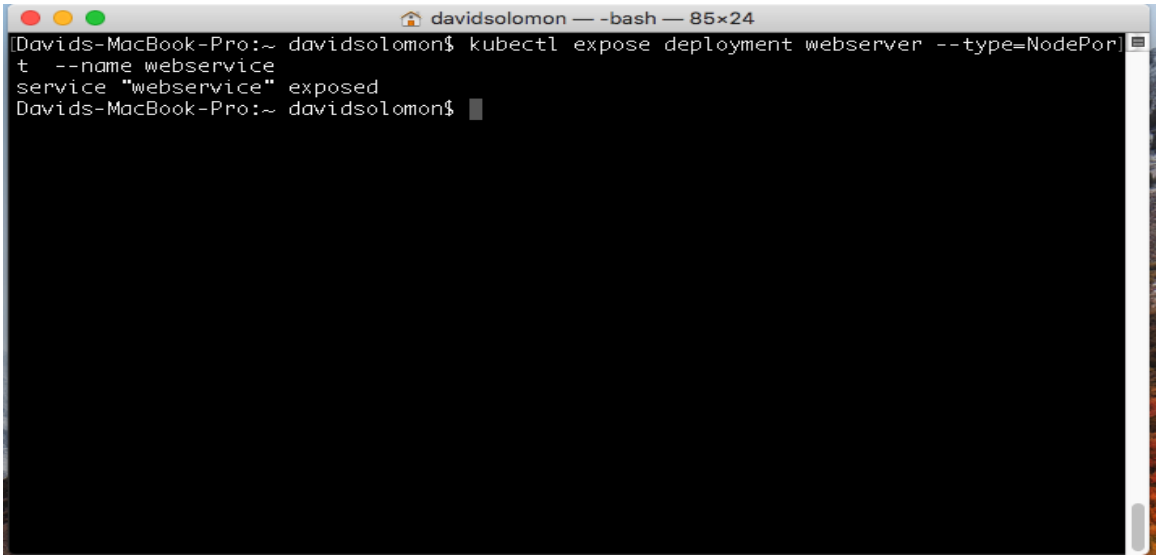 c. Return to the Kubernetes dashboard. You will see that a new service has been created for the application, as shown below, indicating that the application now has a service for accessing it from outside the cluster. 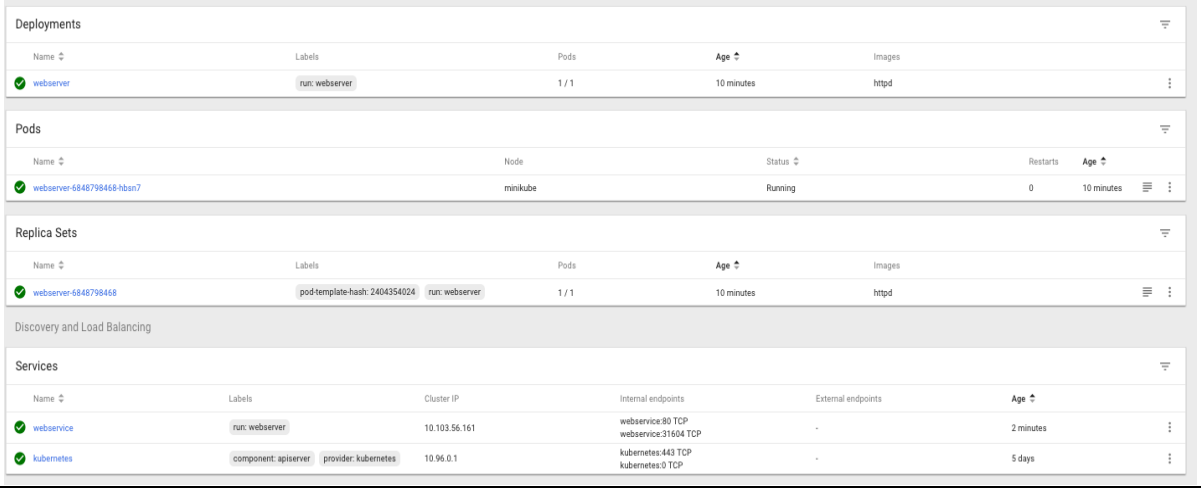

Step	Action
3	**Access the Running Application** a. In order to interact with the application, Kubernetes maintains a set of ports for enabling outside access. These ports are assigned automatically when a service is created and mapped to the port the application is expecting (in this case, port 80). In order to see which port has been assigned for your deployed application, run the following command: ~$ kubectl describe services webservice b. Once this command is entered, you will see the following output. Note the port number listed in the "NodePort" section of the output. This is the port you need to access the application (31604 in the example below). 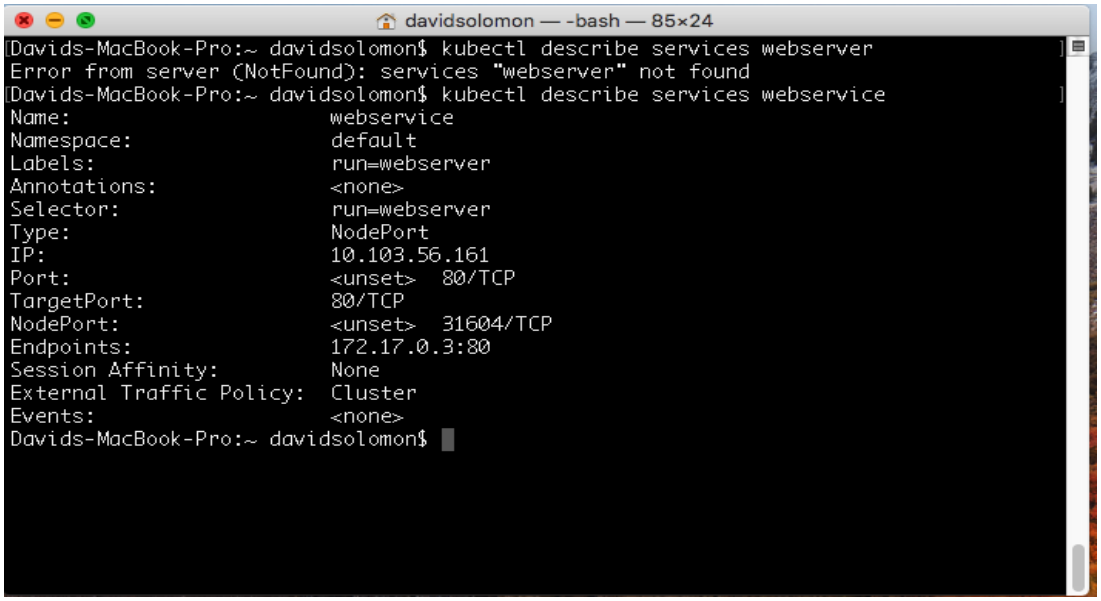 c. To access the application, go to your browser and enter the following URL and verify that you can access the application, as shown below: 192.168.99.100:*<your port number>* 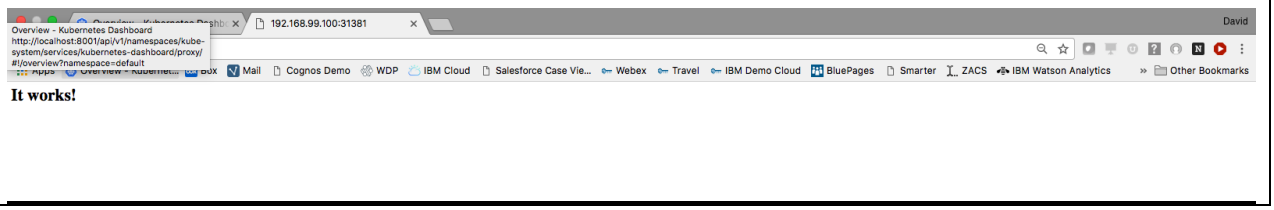

Section 6: Lab Summary

In this section, you learned how to deploy an Docker application to Kubernetes, how to enable it to be access from the outside world, and how to access it.

Section 7: Observing Kubernetes Resiliency

Purpose:	In this lab, you will learn how Kubernetes recovers from a container failure.

Tasks:	Tasks you will complete in this lab exercise include: • Access a container running in Kubernetes from Docker • Simulate a container failure • Observer how the cluster quickly recovers from the failure

Section 7: Observing Kubernetes Resiliency

1. Access a Container Running in Kubernetes
2. Simulate Container Failure
3. Observe How the Cluster Quickly Recovers from Failure

Section 7: Lab Instructions

Step	Action

1	**Access a container running in Kubernetes from Docker**
	
	a. In order to enable the control of containers running inside Kubernetes, we will need to first set your Docker environment so that it can manage these containers. To do this, run the following command:
	~$ minikube docker-env
	
	This command will provide a list of export statements, as shown below:
	
	`Davids-MacBook-Pro:bin davidsolomon$ minikube docker-env`
	`export DOCKER_TLS_VERIFY="1"`
	`export DOCKER_HOST="tcp://192.168.99.100:2376"`
	`export DOCKER_CERT_PATH="/Users/davidsolomon/.minikube/certs"`
	`export DOCKER_API_VERSION="1.23"`
	`# Run this command to configure your shell:`
	`# eval $(minikube docker-env)`
	
	b. Copy the export statements and run them. Confirm that Docker is now pointing to the Kubernetes cluster by entering the following command. Your httpd application should be listed, as shown below:
	~$ docker ps
	
	 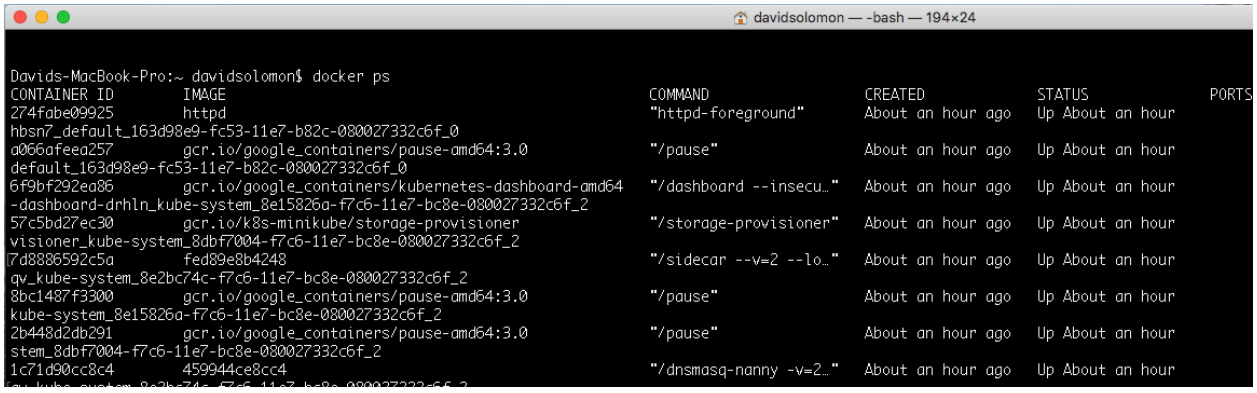

2	**Simulate a Container Failure** a. We will now use a docker command to simulate the failure of the webserver application. To do this, find the container ID for this application from the output shown above, and enter the following command:

Step	Action
	~$ docker kill *<the container ID for your application>* b. Immediately return to and refresh the Kubernetes dashboard. While the container will only take a few seconds to recover, you may see the failure reflected as shown below: 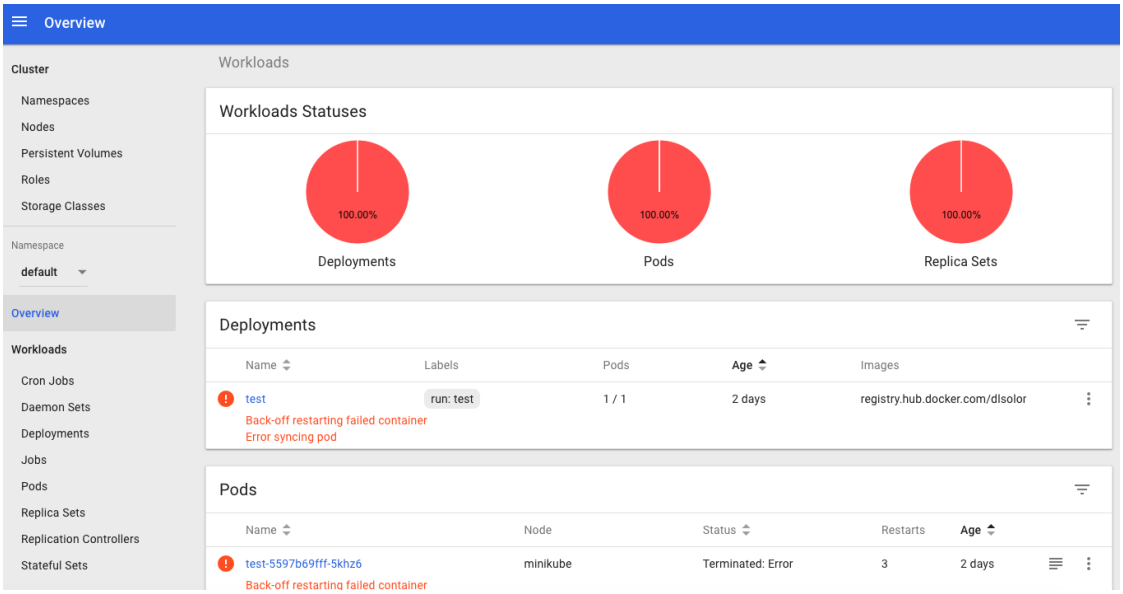
3	**<u>Observe how quickly the cluster recovers from the failure</u>** c. Refresh the dashboard page again. Notice that the status is now green again, and the webserver deployment is now again running.

Step	Action
	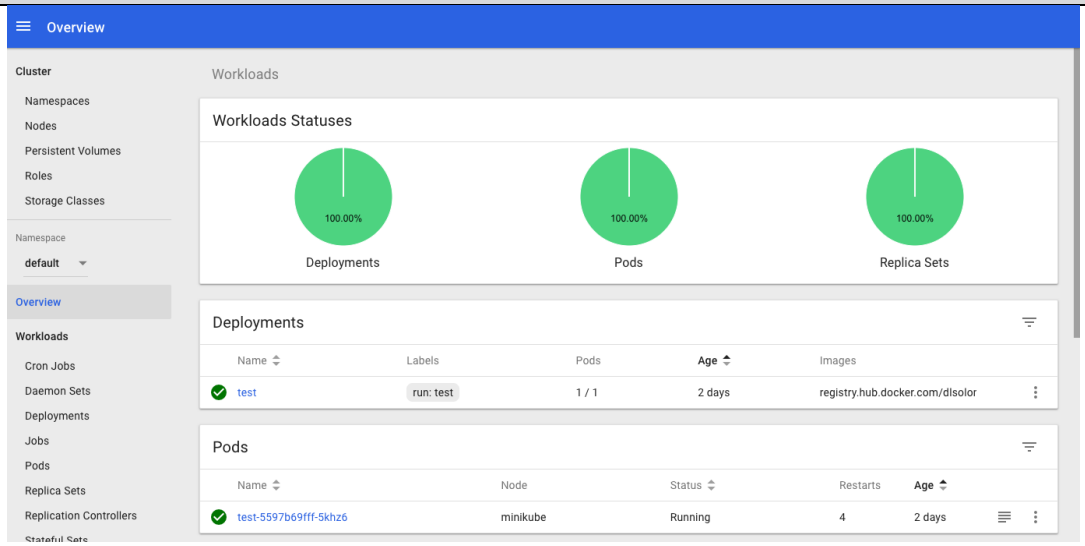 d. We can also observe this behavior by entering the following command and noting that the container ID for “httpd” is now different, indicating that the application has fully recovered as a new container. However, since the exposing service has not changed, there is no change to how the application is accessed: ~$ docker ps

Section 7: Lab Summary

In this section, you learned how Kubernetes can quickly recover from an application failure.